

▼ Bank Customer Churn Prediction

In this project, I use supervised learning models to identify customers who are likely to churn in the factors that influence user retention. [Dataset information](#).

▼ Contents

- [Part 1: Data Exploration](#)
- [Part 2: Feature Preprocessing](#)
- [Part 3: Model Training and Results Evaluation](#)
- [Part 4: Feature Selection](#)

▼ Part 0. oad packages, load data

```
#import neccessary libraries
import numpy as np
import pandas as pd
import sklearn as sl
import sklearn.preprocessing as preprocessing
from sklearn.model_selection import train_test_split
import seaborn as sns
import matplotlib.pyplot as plt
pd.set_option('display.float_format', lambda x: '%.3f' % x)
pd.set_option('display.max_columns',None)
pd.set_option('display.max_rows',None)
pd.set_option('max_colwidth',100)
```

```
↳ /usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarnin
import pandas.util.testing as tm
```

```
from google.colab import files
uploaded = files.upload()
```

```
↳  bank.data.csv
• bank.data.csv(application/vnd.ms-excel) - 684858 bytes, last modified: 3/12/2020 - 100% done
Saving bank.data.csv to bank.data.csv
```

```
churn_df=pd.read_table('bank.data.csv',header=0,sep=',',lineterminator='\n')
print(churn_df.head())
print ("Num of rows: " + str(churn_df.shape[0])) # row count
print ("Num of columns: " + str(churn_df.shape[1])) # col count
```

```

↳
  RowNumber  CustomerId  Surname  CreditScore  Geography  Gender  Age  \
0           1    15634602  Hargrave         619    France  Female  42
1           2    15647311      Hill         608     Spain  Female  41
2           3    15619304      Onio         502    France  Female  42
3           4    15701354     Boni         699    France  Female  39
4           5    15737888  Mitchell         850     Spain  Female  43

```

```

  Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember  \
0       2    0.000              1           1              1
1       1  83807.860              1           0              1
2       8 159660.800              3           1              0
3       1    0.000              2           0              0
4       2 125510.820              1           1              1

```

```

  EstimatedSalary  Exited\r
0      101348.880         1
1      112542.580         0
2      113931.570         1
3       93826.630         0
4       79084.100         0
Num of rows: 10000
Num of columns: 14

```

▼ Part 1. Exploratory Analysis and Data Visualization

▼ 1.1 Exclude erroneous data

```

if churn_df.set_index('CustomerId').index.duplicated().sum()==0:
    print('No duplicated index.')

```

```

↳ No duplicated index.

```

▼ 1.2: Understand the Raw Dataset

```

# check data info
churn_df.info()

```

```

↳

```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber              10000 non-null  int64
1   CustomerId             10000 non-null  int64
```

```
# check the unique values for each column
churn_df.nunique()
```

```
↳ RowNumber      10000
   CustomerId     10000
   Surname        2932
   CreditScore     460
   Geography       3
   Gender          2
   Age            70
   Tenure         11
   Balance       6382
   NumOfProducts  4
   HasCrCard      2
   IsActiveMember 2
   EstimatedSalary 9999
   Exited\r       2
dtype: int64
```

```
# Get target variable
y = churn_df['Exited\r']
```

```
# check the propotion of y = 1
print(y.sum() / y.shape * 100)
```

```
↳ [20.37]
```

▼ 1.3: Understand numerical features

▼ 1.3.1 Overview

```
# check missing values
churn_df.isnull().sum()
```

```
↳
```

```

RowNumber      0
CustomerId      0
Surname         0
CreditScore     0

```

```
churn_df.info()
```

```

↳ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber              10000 non-null  int64
1   CustomerId             10000 non-null  int64
2   Surname                10000 non-null  object
3   CreditScore             10000 non-null  int64
4   Geography              10000 non-null  object
5   Gender                 10000 non-null  object
6   Age                    10000 non-null  int64
7   Tenure                  10000 non-null  int64
8   Balance                 10000 non-null  float64
9   NumOfProducts          10000 non-null  int64
10  HasCrCard               10000 non-null  int64
11  IsActiveMember          10000 non-null  int64
12  EstimatedSalary         10000 non-null  float64
13  Exited                  10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

```

```
print(churn_df.drop(columns=['CustomerId','RowNumber'],axis=1).describe(percentiles=[0.1,0.5,0.9]))
```

```

↳
   CreditScore  Age  Tenure  Balance  NumOfProducts  HasCrCard  \
count  10000.000  10000.000  10000.000  10000.000  10000.000  10000.000
mean     650.529   38.922    5.013  76485.889      1.530      0.706
std      96.653   10.488    2.892  62397.405      0.582      0.456
min     350.000   18.000    0.000    0.000      1.000      0.000
10%     521.000   27.000    1.000    0.000      1.000      0.000
25%     584.000   32.000    3.000    0.000      1.000      0.000
50%     652.000   37.000    5.000  97198.540      1.000      1.000
75%     718.000   44.000    7.000 127644.240      2.000      1.000
95%     812.000   60.000    9.000 162711.669      2.000      1.000
max     850.000   92.000   10.000 250898.090      4.000      1.000

   IsActiveMember  EstimatedSalary  Exited
count  10000.000      10000.000  10000.000
mean      0.515      100090.240    0.204
std      0.500      57510.493    0.403
min      0.000      11.580    0.000
10%      0.000      20273.580    0.000
25%      0.000      51002.110    0.000
50%      1.000      100193.915    0.000
75%      1.000      149388.247    0.000
95%      1.000      190155.375    1.000
max      1.000      199992.480    1.000

```

```
(churn_df==0).sum(axis=0)/churn_df.shape[0]
```

```
↳
```

```

RowNumber      0.000
CustomerId      0.000
Surname         0.000
CreditScore     0.000
Geography      0.000
Gender          0.000
Age            0.000
Tenure         0.041
Balance        0.362
NumOfProducts  0.000
HasCrCard      0.294
IsActiveMember 0.485
EstimatedSalary 0.000
Exited         0.706

```

```

# understand Numerical feature
# discrete/continuous
# 'CreditScore', 'Age', 'Tenure', 'NumberOfProducts'
# 'Balance', 'EstimatedSalary'
churn_df[['CreditScore', 'Age', 'Tenure', 'NumOfProducts', 'Balance', 'EstimatedSalary']].d

```

	CreditScore	Age	Tenure	NumOfProducts	Balance	EstimatedSalary
count	10000.000	10000.000	10000.000	10000.000	10000.000	10000.000
mean	650.529	38.922	5.013	1.530	76485.889	100090.240
std	96.653	10.488	2.892	0.582	62397.405	57510.493
min	350.000	18.000	0.000	1.000	0.000	11.580
25%	584.000	32.000	3.000	1.000	0.000	51002.110
50%	652.000	37.000	5.000	1.000	97198.540	100193.915
75%	718.000	44.000	7.000	2.000	127644.240	149388.247
max	850.000	92.000	10.000	4.000	250898.090	199992.480

```

# Rename columns
churn_df.rename(columns={'Exited\r': 'Exited'}, inplace=True)

```

▼ 1.3.2 'Exited' feature

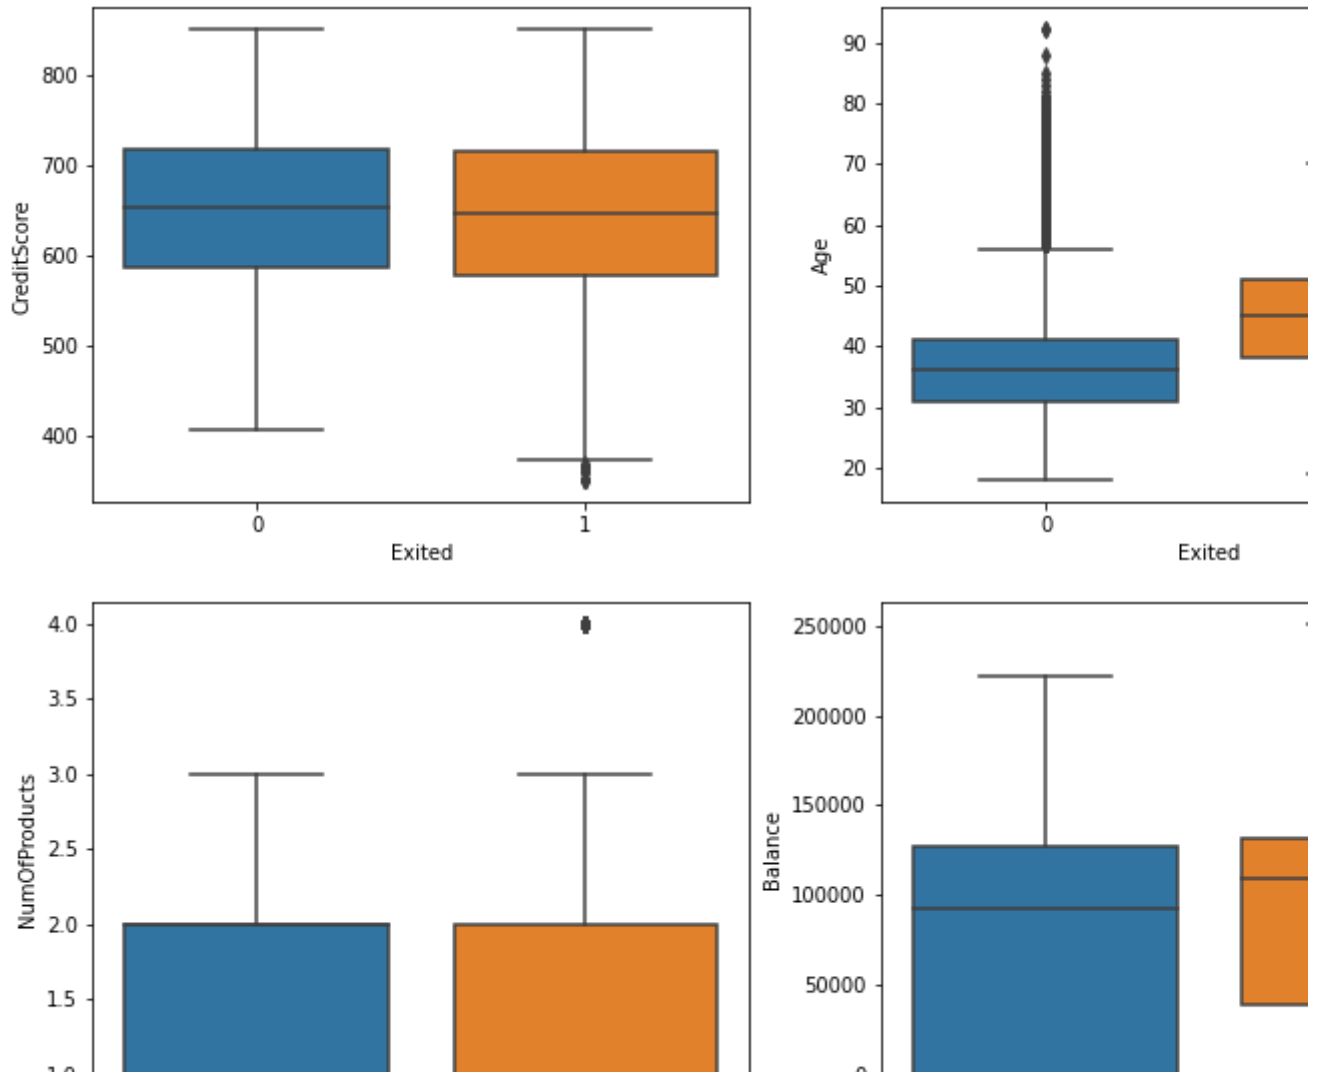
```

# boxplot for numerical feature
_,axss = plt.subplots(2,3, figsize=[20,10])
sns.boxplot(x='Exited', y='CreditScore', data=churn_df, ax=axss[0][0])
sns.boxplot(x='Exited', y='Age', data=churn_df, ax=axss[0][1])
sns.boxplot(x='Exited', y='Tenure', data=churn_df, ax=axss[0][2])
sns.boxplot(x='Exited', y='NumOfProducts', data=churn_df, ax=axss[1][0])
sns.boxplot(x='Exited', y='Balance', data=churn_df, ax=axss[1][1])
sns.boxplot(x='Exited', y='EstimatedSalary', data=churn_df, ax=axss[1][2])

```



<matplotlib.axes._subplots.AxesSubplot at 0x7fb7fc7c0208>



insights:


Senior people are more likely to churn.

People have more balance are more likely to churn.


▼ 1.3.3 Distribution of 'Balance' and 'Estimated Salary'

```
plt.hist(churn_df['Balance'], bins = 10, alpha = 0.4, color='r', histtype='stepfilled', la
plt.legend(loc = 'upper right')
plt.title('Histograms of balance before data processing')
plt.xlabel('blance')
plt.ylabel('Count')
plt.show()
```





```
plt.hist(churn_df['EstimatedSalary'], bins = 500, alpha = 0.4, color='b', histtype='stepfi
plt.legend(loc = 'upper right')
plt.title('Histograms of estimated salary before data processing')
plt.xlabel('estimated salary')
plt.ylabel('Count')
plt.show()
```




▼ 1.2.3 Correlation among numerical features

```
corr_score = churn_df[['CreditScore', 'Age', 'Tenure', 'NumOfProducts', 'Balance', 'Estimat
# show heatmap of correlations
sns.heatmap(corr_score, cmap="YlGnBu")
```



<matplotlib.axes._subplots.AxesSubplot at 0x7fb7fc08edd8>



```
# check the actual values of correlations
corr_score
```



	CreditScore	Age	Tenure	NumOfProducts	Balance	EstimatedSalary
CreditScore	1.000	-0.004	0.001	0.012	0.006	-0.001
Age	-0.004	1.000	-0.010	-0.031	0.028	-0.007
Tenure	0.001	-0.010	1.000	0.013	-0.012	0.008
NumOfProducts	0.012	-0.031	0.013	1.000	-0.304	0.014
Balance	0.006	0.028	-0.012	-0.304	1.000	0.013
EstimatedSalary	-0.001	-0.007	0.008	0.014	0.013	1.000

▼ 1.4 Understand Categorical Features

```
# understand categorical feature
# 'Geography', 'Gender'
# 'HasCrCard', 'IsActiveMember'
_,axss = plt.subplots(2,2, figsize=[20,10])
sns.countplot(x='Exited', hue='Geography', data=churn_df, ax=axss[0][0])
sns.countplot(x='Exited', hue='Gender', data=churn_df, ax=axss[0][1])
sns.countplot(x='Exited', hue='HasCrCard', data=churn_df, ax=axss[1][0])
sns.countplot(x='Exited', hue='IsActiveMember', data=churn_df, ax=axss[1][1])
```



<matplotlib.axes._subplots.AxesSubplot at 0x7fb7f9726940>



Insights:

Female are more likely to churn.

Less active members are more likely to churn.

▼ Part 2: Feature Preprocessing

feature encoding, feature scaling

After very basic Exploratory Data Analysis, we have to do some data cleaning and data preprocess we need to encode the categorical feature Second, we need to impute the missing value for both n need to scale out feature,which can be better for our models' performance

Read more for handling [categorical feature](#), and there is an awesome package for [encoding](#).

Exited

```
# ordinal encoding
```

```
churn_df['Gender'] = churn_df['Gender'] == 'Female'
```

```
# one hot encoding
```

```
churn_df = pd.get_dummies(churn_df, columns=['Geography'], drop_first=True)
```

```
churn_df.head(10)
```



	RowNumber	CustomerId	Surname	CreditScore	Gender	Age	Tenure	Balance	NumO
0	1	15634602	Hargrave	619	True	42	2	0.000	
1	2	15647311	Hill	608	True	41	1	83807.860	

```
# Get feature space by dropping useless feature
to_drop = ['RowNumber', 'CustomerId', 'Surname', 'Exited']
X = churn_df.drop(to_drop, axis=1)
```

```
X.head()
```

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMer
0	619	True	42	2	0.000	1	1	
1	608	True	41	1	83807.860	1	0	
2	502	True	42	8	159660.800	3	1	
3	699	True	39	1	0.000	2	0	
4	850	True	43	2	125510.820	1	1	

▼ Part 3: Model Training and Result Evaluation

▼ 3.1 Feature scaling

The impact of different scaling methods on the model performance is small. In the following model, standardization (sc) data is used.

```
# Scale the data, using standardization
# standardization (x-mean)/std
# normalization (x-x_min)/(x_max-x_min) ->[0,1]

# 1. speed up gradient descent
# 2. same scale
# 3. algorithm requirments

# for example, use training data to train the standardscaler to get mean and std
# apply mean and std to both training and testing data.
# fit_transform does the training and applying, transform only does applying.
# Because we can't use any info from test, and we need to do the same modification
# to testing data as well as training data

# https://scikit-learn.org/stable/auto\_examples/preprocessing/plot\_all\_scaling.html#sphx-g
# https://scikit-learn.org/stable/modules/preprocessing.html
```

```
# min_max example: (x - x_min) / (x_max - x_min)
```

```
# min-max example. (x-x_min)/(x_max-x_min)
# [1,2,3,4,5,6] -> fit(min:1, max:6) (scalar.min = 1, scalar.max = 6) -> transform [(1-1)/
# scalar.fit(train) -> min:1, max:100
# scalar.transform(apply to x) -> apply min:1, max:100 to X_train
# scalar.transform -> apply min:1, max:100 to X_test

# scalar.fit -> mean:1, std:100
# scalar.transform -> apply mean:1, std:100 to X_train
# scalar.transform -> apply mean:1, std:100 to X_test

from sklearn.preprocessing import StandardScaler
scale_lst = ['CreditScore', 'Age', 'Tenure', 'NumOfProducts', 'Balance', 'EstimatedSalary']
scaler = StandardScaler()
scaler.fit(X[scale_lst])
X[scale_lst] = scaler.transform(X[scale_lst])

X.head()
```

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMeml
0	-0.326	True	0.294	-1.042	-1.226	-0.912	1	
1	-0.440	True	0.198	-1.388	0.117	-0.912	0	
2	-1.537	True	0.294	1.033	1.333	2.527	1	
3	0.502	True	0.007	-1.388	-1.226	0.808	0	
4	2.064	True	0.389	-1.042	0.786	-0.912	1	

▼ 3.2 Split the data

```
# Splite data into training and testing
from sklearn import model_selection

# Reserve 20% for testing
X_train, X_test, y_train, y_test = model_selection.train_test_split(X, y, test_size=0.25,

print('training data has ' + str(X_train.shape[0]) + ' observation with ' + str(X_train.sh
print('test data has ' + str(X_test.shape[0]) + ' observation with ' + str(X_test.shape[1]

training data has 7500 observation with 11 features
test data has 2500 observation with 11 features
```

▼ 3.3: Model Training and Selection

```
# build models
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn import svm
```

```
# Logistic Regression
```

```

# Logistic Regression
classifier_logistic = LogisticRegression()

# K Nearest Neighbors
classifier_KNN = KNeighborsClassifier()

# Random Forest
classifier_RF = RandomForestClassifier()

#SVM
classifier_SVM = svm.SVC()

# Train the model
#classifier_logistic.fit(X_train, y_train)

# Prediction of test data
#classifier_logistic.predict(X_test)

# Accuracy of test data
#classifier_logistic.score(X_test, y_test)

# Use 5-fold Cross Validation to get the accuracy for different models
model_names = ['Logistic Regression','KNN','Random Forest','SVM']
model_list = [classifier_logistic, classifier_KNN, classifier_RF, classifier_SVM]
count = 0

for classifier in model_list:
    if count<=len(model_list):
        cv_score = model_selection.cross_val_score(classifier, X_train, y_train, cv=5)
        print(cv_score)
        print('Model accuracy of ' + model_names[count] + ' is ' + str(cv_score.mean()))
        count += 1

[0.81      0.81466667 0.802      0.818      0.81266667]
Model accuracy of Logistic Regression is 0.8114666666666667
[0.852      0.83333333 0.846      0.85133333 0.848      ]
Model accuracy of KNN is 0.8461333333333334
[0.86133333 0.868      0.86333333 0.86533333 0.862      ]
Model accuracy of Random Forest is 0.8640000000000001
[0.86      0.86266667 0.85133333 0.85866667 0.85666667]
Model accuracy of SVM is 0.8578666666666667

```

▼ 3.4 Use Grid Search to Find Optimal Hyperparameters

```

from sklearn.model_selection import GridSearchCV

# helper function for printing out grid search results
def print_grid_search_metrics(gs):
    print ("Best score: " + str(gs.best_score_))
    print ("Best parameters set:")
    best_parameters = gs.best_params_
    for param_name in sorted(parameters.keys()):

```

```

for param_name in sorted(parameters.keys()):
    print(param_name + ':' + str(best_parameters[param_name]))

```

▼ 3.4.1 Find Optimal Hyperparameters - LogisticRegression

```

# Possible hyperparameter options for Logistic Regression Regularization
# Penalty is choosed from L1 or L2
# C is the lambda value(weight) for L1 and L2

# ('l1', 1) ('l1', 5) ('l1', 10) ('l2', 1) ('l2', 5) ('l2', 10)
parameters = {
    'penalty':('l1', 'l2'),
    'C':(0.01, 0.1, 1, 5, 10)
}
Grid_LR = GridSearchCV(LogisticRegression(solver='liblinear'),parameters, cv=5)
Grid_LR.fit(X_train, y_train)

↳ GridSearchCV(cv=5, error_score=nan,
                estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                                              fit_intercept=True,
                                              intercept_scaling=1, l1_ratio=None,
                                              max_iter=100, multi_class='auto',
                                              n_jobs=None, penalty='l2',
                                              random_state=None, solver='liblinear',
                                              tol=0.0001, verbose=0,
                                              warm_start=False),
                iid='deprecated', n_jobs=None,
                param_grid={'C': (0.01, 0.1, 1, 5, 10), 'penalty': ('l1', 'l2')},
                pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                scoring=None, verbose=0)

# the best hyperparameter combination
print_grid_search_metrics(Grid_LR)

↳ Best score: 0.8134666666666666
   Best parameters set:
   C:0.1
   penalty:l1

# best model
best_LR_model = Grid_LR.best_estimator_

```

▼ 3.4.2 Find Optimal Hyperparameters: KNN

```

# Possible hyperparameter options for KNN
# Choose k
parameters = {
    'n_neighbors':[1,3,5,7,9]
}
Grid_KNN = GridSearchCV(KNeighborsClassifier(),parameters, cv=5)
Grid_KNN.fit(X_train, y_train)

```

```

GridSearchCV(cv=5, error_score=nan,
             estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30,
                                           metric='minkowski',
                                           metric_params=None, n_jobs=None,
                                           n_neighbors=5, p=2,
                                           weights='uniform'),
             iid='deprecated', n_jobs=None,
             param_grid={'n_neighbors': [1, 3, 5, 7, 9]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring=None, verbose=0)

```

```

# best k
print_grid_search_metrics(Grid_KNN)

```

```

Best score: 0.8492000000000001
Best parameters set:
n_neighbors:9

```

```
best_KNN_model = Grid_KNN.best_estimator_
```

▼ 3.4.3 Find Optimal Hyperparameters: Random Forest

```

# Possible hyperparameter options for Random Forest
# Choose the number of trees
parameters = {
    'n_estimators' : [40,60,80]
}
Grid_RF = GridSearchCV(RandomForestClassifier(),parameters, cv=5)
Grid_RF.fit(X_train, y_train)

```

```

GridSearchCV(cv=5, error_score=nan,
             estimator=RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
                                           class_weight=None,
                                           criterion='gini', max_depth=None,
                                           max_features='auto',
                                           max_leaf_nodes=None,
                                           max_samples=None,
                                           min_impurity_decrease=0.0,
                                           min_impurity_split=None,
                                           min_samples_leaf=1,
                                           min_samples_split=2,
                                           min_weight_fraction_leaf=0.0,
                                           n_estimators=100, n_jobs=None,
                                           oob_score=False,
                                           random_state=None, verbose=0,
                                           warm_start=False),
             iid='deprecated', n_jobs=None,
             param_grid={'n_estimators': [40, 60, 80]}, pre_dispatch='2*n_jobs',
             refit=True, return_train_score=False, scoring=None, verbose=0)

```

```

# best number of trees
print_grid_search_metrics(Grid_RF)

```

```


```

```
Best score: 0.8614666666666666
```

```
Best parameters set:
```

```
# best random forest
```

```
best_RF_model = Grid_RF.best_estimator_
```

▼ 3.4.4 Find Optimal Hyperparameters: Support Vector Machines

```
# Possible hyperparameter options for SVM
```

```
# Choose kernel
```

```
parameters = {'kernel': ['rbf'], 'gamma': [1e-3, 1e-4],  
              'C': [1, 10, 100, 1000]}
```

```
Grid_SVM = GridSearchCV(svm.SVC(), parameters, cv=5)
```

```
Grid_SVM.fit(X_train, y_train)
```

```
↳ GridSearchCV(cv=5, error_score=nan,  
               estimator=SVC(C=1.0, break_ties=False, cache_size=200,  
                             class_weight=None, coef0=0.0,  
                             decision_function_shape='ovr', degree=3,  
                             gamma='scale', kernel='rbf', max_iter=-1,  
                             probability=False, random_state=None, shrinking=True,  
                             tol=0.001, verbose=False),  
               iid='deprecated', n_jobs=None,  
               param_grid={'C': [1, 10, 100, 1000], 'gamma': [0.001, 0.0001],  
                           'kernel': ['rbf']},  
               pre_dispatch='2*n_jobs', refit=True, return_train_score=False,  
               scoring=None, verbose=0)
```

```
# best gamma and C
```

```
print_grid_search_metrics(Grid_SVM)
```

```
↳ Best score: 0.8512000000000001  
Best parameters set:  
C:1000  
gamma:0.001  
kernel:rbf
```

```
# best SVM model
```

```
best_SVM_model = Grid_SVM.best_estimator_
```

▼ 3.5 Model Evaluation

class of interest as positive

TP: correctly labeled real churn

Precision (PPV, positive predictive value): $tp / (tp + fp)$; Total number of true predictive churn divide
High Precision means low fp, not many return users were predicted as churn users.

Recall (sensitivity, hit rate, true positive rate): $tp / (tp + fn)$ Predict most positive or churn user correct
users were predicted as return users.

▼ 3.5.1 Model Evaluation - Confusion Matrix (Precision, Recall, Accuracy)

```

from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score

# calculate accuracy, precision and recall, [[tn, fp],[fn, tp]]
def cal_evaluation(classifier, cm):
    tn = cm[0][0]
    fp = cm[0][1]
    fn = cm[1][0]
    tp = cm[1][1]
    accuracy = (tp + tn) / (tp + fp + fn + tn + 0.0)
    precision = tp / (tp + fp + 0.0)
    recall = tp / (tp + fn + 0.0)
    print(classifier)
    print("Accuracy is: " + str(accuracy))
    print("precision is: " + str(precision))
    print("recall is: " + str(recall))

# print out confusion matrices
def draw_confusion_matrices(confusion_matrices):
    class_names = ['Retain', 'Churn']
    for cm in confusion_matrices:
        classifier, cm = cm[0], cm[1]
        cal_evaluation(classifier, cm)
        fig = plt.figure()
        ax = fig.add_subplot(111)
        cax = ax.matshow(cm, interpolation='nearest', cmap=plt.get_cmap('Reds'))
        plt.title('Confusion matrix for ' + classifier)
        fig.colorbar(cax)
        ax.set_xticklabels([''] + class_names)
        ax.set_yticklabels([''] + class_names)
        plt.xlabel('Predicted')
        plt.ylabel('True')
        plt.show()

#confusion_matrix(y_test,best_RF_model.predict(X_test))

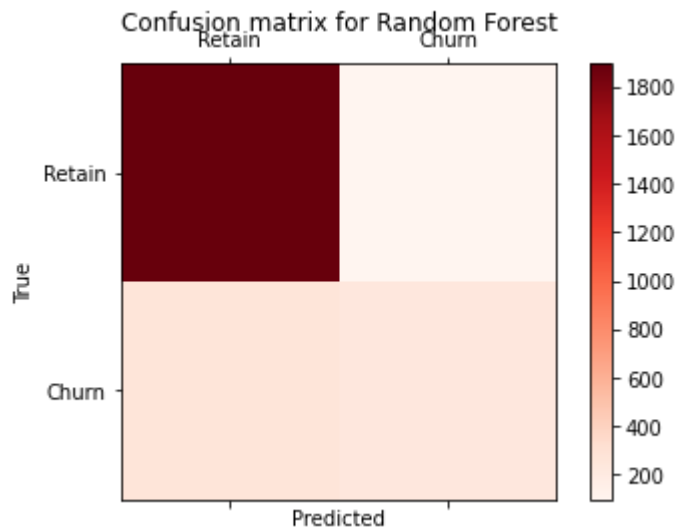
# Confusion matrix, accuracy, precison and recall for random forest and logistic regressio
confusion_matrices = [
    ("Random Forest", confusion_matrix(y_test,best_RF_model.predict(X_test))),
    ("Logistic Regression", confusion_matrix(y_test,best_LR_model.predict(X_test))),
    ("K nearest neighbor", confusion_matrix(y_test, best_KNN_model.predict(X_test))),
    ('Support Vector Moachines', confusion_matrix(y_test, best_SVM_model.predict(X_test)))
]

draw_confusion_matrices(confusion_matrices)


```



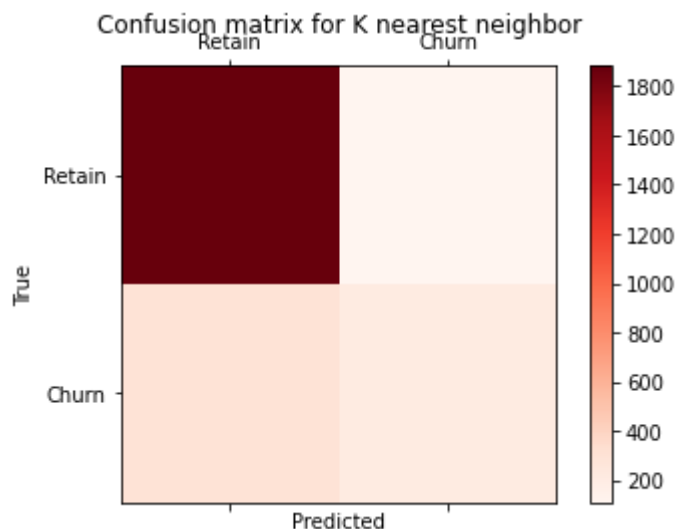
Random Forest
Accuracy is: 0.8536
precision is: 0.7109144542772862
recall is: 0.47347740667976423



Logistic Regression
Accuracy is: 0.8044
precision is: 0.5515463917525774
recall is: 0.21021611001964635




K nearest neighbor
Accuracy is: 0.838
precision is: 0.6614906832298136
recall is: 0.41846758349705304



Support Vector Moachines

Accuracy is: 0.8576
 precision is: 0.8558139534883721
 recall is: 0.3614931237721022



Random Forest has the best performance

▼ 3.5.2 Model Evaluation - ROC & AUC


ROC of RF Model

```
from sklearn.metrics import roc_curve
from sklearn import metrics

# Use predict_proba to get the probability results of Random Forest
y_pred_rf = best_RF_model.predict_proba(X_test)[: , 1]
fpr_rf, tpr_rf, _ = roc_curve(y_test, y_pred_rf)

# ROC curve of Random Forest result
plt.figure(1)
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_rf, tpr_rf, label='RF')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.title('ROC curve - RF model')
plt.legend(loc='best')
plt.show()
```





```
from sklearn import metrics
```

```
# AUC score
```

```
metrics.auc(fpr_rf, tpr_rf)
```

```
0.8424501612857072
```

ROC of LR Model

False positive rate

```
# Use predict_proba to get the probability results of Logistic Regression
```

```
y_pred_lr = best_LR_model.predict_proba(X_test)[: , 1]
```

```
fpr_lr, tpr_lr, thres = roc_curve(y_test, y_pred_lr)
```

```
# ROC Curve
```

```
plt.figure(1)
```

```
plt.plot([0, 1], [0, 1], 'k--')
```

```
plt.plot(fpr_lr, tpr_lr, label='LR')
```

```
plt.xlabel('False positive rate')
```


```
plt.ylabel('True positive rate')
```

```
plt.title('ROC curve - LR Model')
```

```
plt.legend(loc='best')
```

```
plt.show()
```

```
0.745923453181754
```



```
# AUC score
```

```
metrics.auc(fpr_lr, tpr_lr)
```


```
0.745923453181754
```

ROC of KNN Model

```
y_pred_knn = best_KNN_model.predict(X_test)
```

```
y_pred_knn = best_knn_model.predict(X_test)
fpr_knn, tpr_knn, thres = roc_curve(y_test, y_pred_knn)
```

```
# ROC Curve
plt.figure(1)
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_knn, tpr_knn, label='LR')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.title('ROC curve - KNN Model')
plt.legend(loc='best')
plt.show()
```



```
metrics.auc(fpr_knn, tpr_knn)
```



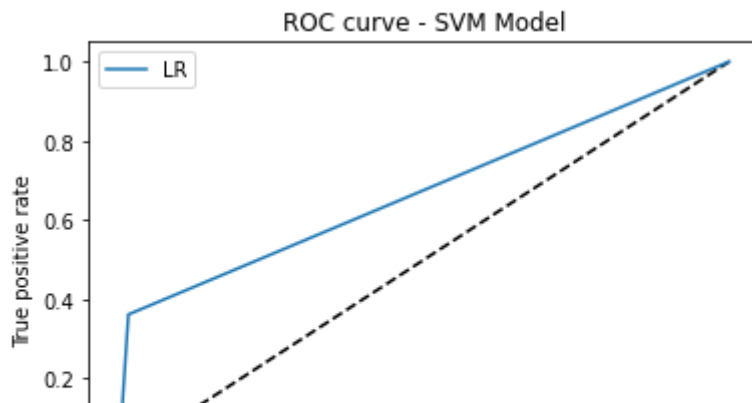
```
0.6818606124416455
```

ROC of SVM Model

```
y_pred_svm = best_SVM_model.predict(X_test)
fpr_svm, tpr_svm, thres = roc_curve(y_test, y_pred_svm)
```

```
# ROC Curve
plt.figure(1)
plt.plot([0, 1], [0, 1], 'k--')
plt.plot(fpr_svm, tpr_svm, label='LR')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.title('ROC curve - SVM Model')
plt.legend(loc='best')
plt.show()
```





```
metrics.auc(fpr_svm, tpr_svm)
```

```
0.67296152923914
```

▼ Part 4: Feature Importance

▼ 4.1 Logistic Regression Model - Feature Selection Discussion

The correlated features that we are interested in

```
# add L1 regularization to logistic regression
# check the coef for feature selection
scaler = StandardScaler()
X_l1 = scaler.fit_transform(X)
LRmodel_l1 = LogisticRegression(penalty="l1", C = 0.01, solver='liblinear')
LRmodel_l1.fit(X_l1, y)

indices = np.argsort(abs(LRmodel_l1.coef_[0]))[::-1]

print ("Logistic Regression (L1) Coefficients")
for ind in range(X.shape[1]):
    print ("{0} : {1}".format(X.columns[indices[ind]], round(LRmodel_l1.coef_[0][indices[ind]]
```

```
Logistic Regression (L1) Coefficients
```


```
Age : 0.6469
IsActiveMember : -0.419
Geography_Germany : 0.2719
Gender : 0.1802
Balance : 0.1133
NumOfProducts : -0.0029
Geography_Spain : 0.0
EstimatedSalary : 0.0
HasCrCard : 0.0
Tenure : 0.0
CreditScore : 0.0
```

```
feature_name = X_train.columns.values
importances_l1 = abs(LRmodel_l1.coef_[0])
plt.figure(1)
```

```

#----->
axes = plt.gca()
plt.bar(feature_name[indices], importances_l1[indices])
plt.title('feature importance of L1 model')
plt.xticks(rotation=90)
#axes.set_ylim([-0.2,0.5])
plt.show()

```



```

# add L2 regularization to logistic regression
# check the coef for feature selection
np.random.seed()
scaler = StandardScaler()
X_l2 = scaler.fit_transform(X)
LRmodel_l2 = LogisticRegression(penalty="l2", C = 0.1, solver='liblinear', random_state=42)
LRmodel_l2.fit(X_l2, y)
LRmodel_l2.coef_[0]

```

```
indices = np.argsort(abs(LRmodel_l2.coef_[0]))[::-1]
```

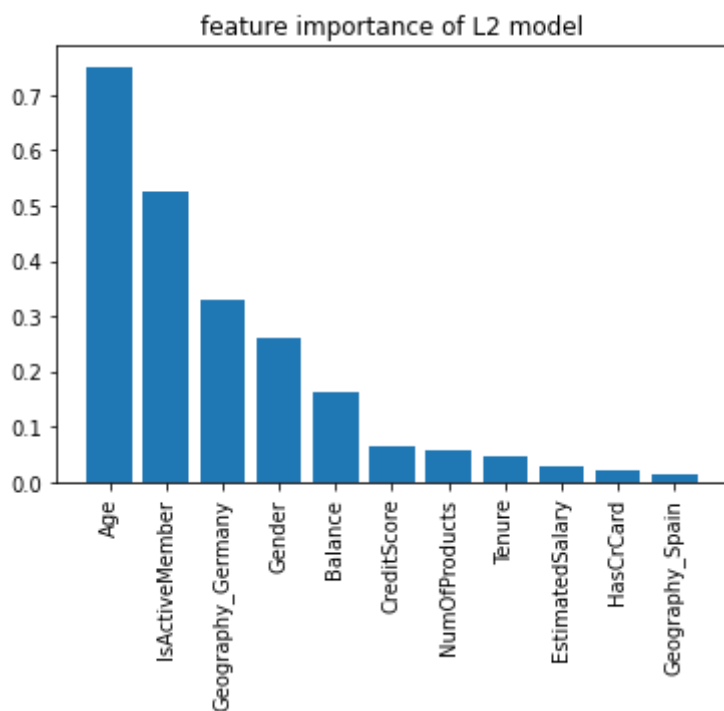
```

print ("Logistic Regression (L2) Coefficients")
for ind in range(X.shape[1]):
    print ("{0} : {1}".format(X.columns[indices[ind]],round(LRmodel_l2.coef_[0][indices[ind]]

```



```
feature_name = X_train.columns.values
plt.figure(1)
importances_l2 = abs(LRmodel_l2.coef_[0])
axes = plt.gca()
plt.bar(feature_name[indices], importances_l2[indices])
plt.title('feature importance of L2 model')
plt.xticks(rotation=90)
plt.show()
```



▼ 4.2 Random Forest Model - Feature Importance Discussion

```
# check feature importance of random forest for feature selection
forest = RandomForestClassifier()
forest.fit(X, y)

importances = forest.feature_importances_

indices = np.argsort(importances)[::-1]

# Print the feature ranking
print("Feature importance ranking by Random Forest Model:")
for ind in range(X.shape[1]):
    print("{0} : {1}".format(X.columns[indices[ind]], round(importances[indices[ind]], 4)))
```



Feature importance ranking by Random Forest Model:

Age : 0.2381

EstimatedSalary : 0.1458

CreditScore : 0.1449

```
feature_name = X_train.columns.values
plt.figure(1)
plt.bar(feature_name[indices], importances[indices])
plt.title('feature importance of Random Forest model')
plt.xticks(rotation=90)
plt.show()
```

